

# Introduction to Statistical Learning


Omid Safarzadeh

January 19, 2022

# Table of contents

- 1 Sequence to sequence models
- 2 Attention Mechanism
  - Bottleneck Problem
  - Attention Layer
- 3 Categories of Attention Mechanism
- 4 Sequence to sequence models
- 5 Transformers
  - Self attention mechanism
  - Multi-Head attention mechanism
  - Encoder Architecture
  - Decoder Architecture
  - Full Architecture
- 6 Positional Encoding
- 7 BERT
  - Language Masked Learning
  - BERT Input
  - BERT Output

# RNN Structure



This structure allows RNNs to link information from earlier steps to a present step. However, if the RNN is very deep, meaning that it has many layers, it will be prone to vanishing or exploding gradients.

# RNN model

Given a sequence of  $X = (x_1, x_2, \dots, x_T)$  inputs, RNN estimates a sequence of  $Y = (y_1, y_2, \dots, y_T)$  outputs by iterating the following equation:

$$h_t = \text{Sigmoid}(W^{hx}x_t + W^{hh}h_{t-1})$$
$$y_t = W^{yh}h_t$$

Problem?  $T$  should be fixed in both input and output sequences. Solution? map input sequence to a fixed length vector and then map that vector to target sequence with another RNN.

# Encoder - Decoder Architecture

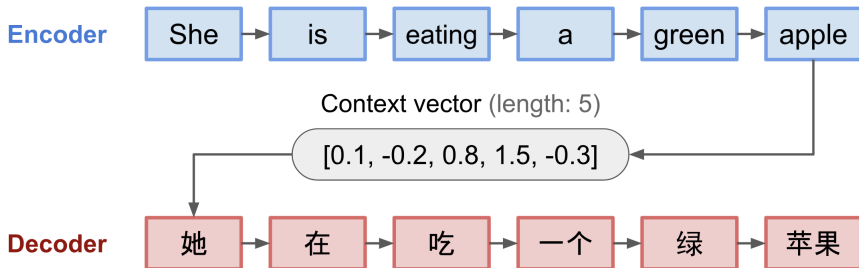
- proposed by (Cho et. al., 2014)
- The **encoder** processes each item in the input sequence, it compiles the information it captures into **context vector**. After processing the entire input sequence, the encoder sends the context vector to the **decoder**, which begins producing the output sequence item by item.

Example Example

# Sequence to Sequence models

- A Sequence to Sequence model,, is a model that takes a sequence of items (words, letters, features , ...) and outputs another sequence of items (with any length for both input and output).
- Applications : machine translation between multiple languages in either text or audio, question-answer dialog generation, or even parsing sentences into grammar trees.
- seq2seq is introduced by (Sutskever, et al. 2014) replaced RNN with LSTM models. **Example**

# NMT Architecture



Use RNN networks (LSTM, GRU,...) for Encoder and Decoder sides. Also we need a Context Vector. (here fixed length).

Context Vector is also called sentence embedding or "thought" vector which represents a good summary of the meaning of the whole source sequence.

# Attention Mechanism



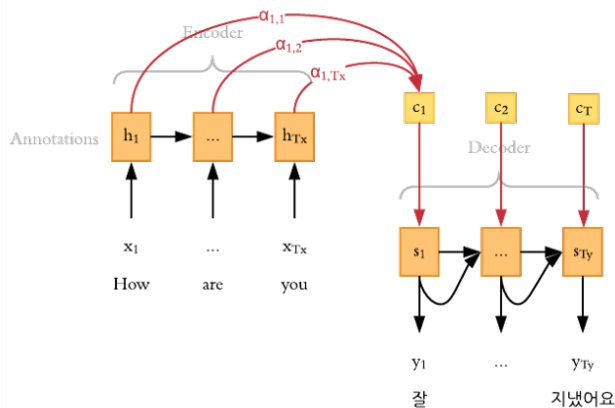
## "Definition "

attention can be interpreted as a vector of importance weights: in order to predict or infer one element, such as a pixel in an image or a word in a sentence, we estimate using the attention vector how strongly it is correlated with ( "attends to" ) other elements and take the sum of their values weighted by the attention vector as the approximation of the target.

Ref: Lilian Weng

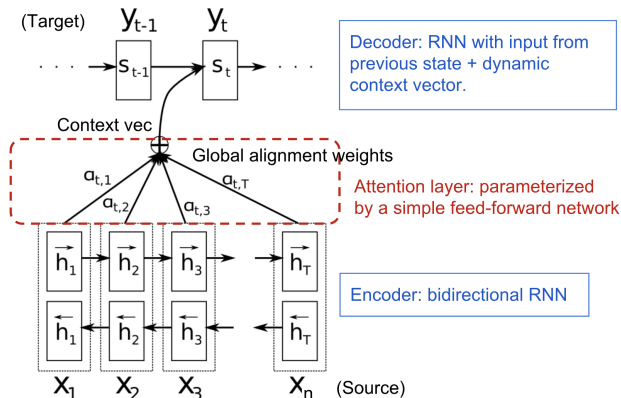
# Bottleneck problem

The last hidden state is actually the context we pass along to the decoder and it is big challenge for the models to deal with long sentences.



# Attention

Attention allows the model to focus on the relevant parts of the input sequence as needed. First, the encoder passes a lot more data to the decoder. Instead of passing the last hidden state of the encoding stage, the encoder passes all the hidden states to the decoder:



**Additive Attention**

Bahdanau, et al. (2014).

$$\begin{aligned}p(y) &= \prod_t^T x(p(y_t|y_1, \dots, y_{t-1}), c) \\p((y_t|y_1, \dots, y_{t-1}), c) &= g(y_{t-1}, s_t, c) \\s_t &= f(s_{t-1}, y_{t-1}, c_t) \\c_t &= \sum_j^T x(\alpha_{tj} h_j) \\\alpha_{tj} &= \frac{\exp(e_{tj})}{\sum_k^T x(e_{tk})} \\e_{tj} &= a(s_{t-1}, h_j)\end{aligned}$$

C: context vector of output  $y_t$   
h: hidden state ;  $\alpha_{tj}$ : Softmax

First, the encoder passes a lot more data to the decoder. Instead of passing the last hidden state of the encoding stage, the encoder passes all the hidden states to the decoder [Example](#)

Second : the decoder does the following:

- Look at the set of encoder hidden states it received – each encoder hidden states is most associated with a certain word in the input sentence
- Give each hidden states a score (lets ignore how the scoring is done for now)
- Multiply each hidden states by its softmaxed score, thus amplifying hidden states with high scores, and drowning out hidden states with low scores

Example

# Attention - all steps:

- 1 The attention decoder RNN takes in the embedding of the END token, and an initial decoder hidden state.
- 2 The RNN processes its inputs, producing an output and a new hidden state vector ( $h_4$ ). The output is discarded.
- 3 Attention Step: We use the encoder hidden states and the  $h_4$  vector to calculate a context vector ( $C_4$ ) for this time step.
- 4 We concatenate  $h_4$  and  $C_4$  into one vector.
- 5 We pass this vector through a feedforward neural network (one trained jointly with the model).
- 6 The output of the feedforward neural networks indicates the output word of this time step.
- 7 Repeat for the next time steps

## Example

# Categories of Attention Mechanisms

- Self – Attention
- Local vs Global Attention
- Hard vs Soft Attention



# Self – attention (intra-attention)

- Machine reading
- Abstractive Summarization
- Image description generation

Ref:Loung 2015

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .


The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

# Image Captioning with Attention




A woman is throwing a frisbee in a park.



A dog is standing on a hardwood floor.




A stop sign is  
mountain in t



A little girl sitting on a bed with  
a teddy bear.



A group of people sitting on a boat  
in the water.




A giraffe standi  
trees in the bac

# Transformers

# Self-attention mechanism

Embedding size: 512 ,  $q_1$  (1\*64) , key(1\*64) , value(1\*64)



# Matrix Notation

$W^Q$ ,  $W^K$ ,  $W^V$  : have **512\*64** weights to learn.

X dimensions is (**maxlen \* 512**). So X represents a sentence with two words.

$$\begin{matrix} \text{X} \\ \begin{array}{|c|c|c|c|} \hline \square & \square & \square & \square \\ \hline \square & \square & \square & \square \\ \hline \end{array} \end{matrix} \times \begin{matrix} W^Q \\ \begin{array}{|c|c|c|c|} \hline \square & \square & \square & \square \\ \hline \square & \square & \square & \square \\ \hline \end{array} \end{matrix} = \begin{matrix} Q \\ \begin{array}{|c|c|} \hline \square & \square \\ \hline \square & \square \\ \hline \end{array} \end{matrix}$$

$$\begin{matrix} \text{X} \\ \begin{array}{|c|c|c|c|} \hline \square & \square & \square & \square \\ \hline \square & \square & \square & \square \\ \hline \end{array} \end{matrix} \times \begin{matrix} W^K \\ \begin{array}{|c|c|c|c|} \hline \square & \square & \square & \square \\ \hline \square & \square & \square & \square \\ \hline \end{array} \end{matrix} = \begin{matrix} K \\ \begin{array}{|c|c|} \hline \square & \square \\ \hline \square & \square \\ \hline \end{array} \end{matrix}$$

$$\begin{matrix} \text{X} \\ \begin{array}{|c|c|c|c|} \hline \square & \square & \square & \square \\ \hline \square & \square & \square & \square \\ \hline \end{array} \end{matrix} \times \begin{matrix} W^V \\ \begin{array}{|c|c|c|c|} \hline \square & \square & \square & \square \\ \hline \square & \square & \square & \square \\ \hline \end{array} \end{matrix} = \begin{matrix} V \\ \begin{array}{|c|c|} \hline \square & \square \\ \hline \square & \square \\ \hline \end{array} \end{matrix}$$

# Calculate Self Attention Score

- $q_1 * k_1$  score for position 1
- $q_1 * k_2$  score for position 2
- $d_k = 64$

$$\text{softmax}\left(\frac{\begin{matrix} \text{Q} \\ \begin{matrix} \square & \square & \square \end{matrix} \end{matrix} \times \begin{matrix} \text{K}^T \\ \begin{matrix} \square \\ \square \\ \square \end{matrix} \end{matrix}}{\sqrt{d_k}}\right) \begin{matrix} \text{V} \\ \begin{matrix} \square & \square & \square \end{matrix} \end{matrix}$$
$$= \begin{matrix} \text{Z} \\ \begin{matrix} \square & \square & \square \end{matrix} \end{matrix}$$

# Tensorflow implementation


## use tf.linalg.matmul for tensorflow 2

```
def scaled_dot_product_attention(q, k, v, mask):  
    """Calculate the attention weights.  
    q, k, v must have matching leading dimensions.  
    k, v must have matching penultimate dimension, i.e.: seq_len_k = seq_len_v.  
    The mask has different shapes depending on its type(padding or look ahead)  
    but it must be broadcastable for addition.  
  
    Args:  
        q: query shape == (... , seq_len_q, depth)  
        k: key shape == (... , seq_len_k, depth)  
        v: value shape == (... , seq_len_v, depth_v)  
        mask: Float tensor with shape broadcastable  
              to (... , seq_len_q, seq_len_k). Defaults to None.  
  
    Returns:  
        output, attention_weights  
    """  
  
    matmul_qk = tf.matmul(q, k, transpose_b=True) # (... , seq_len_q, seq_len_k)  
  
    # scale matmul_qk  
    dk = tf.cast(tf.shape(k)[-1], tf.float32)  
    scaled_attention_logits = matmul_qk / tf.math.sqrt(dk)  
  
    # add the mask to the scaled tensor.  
    if mask is not None:  
        scaled_attention_logits += (mask * -1e9)  
  
    # softmax is normalized on the last axis (seq_len_k) so that the scores  
    # add up to 1.  
    attention_weights = tf.nn.softmax(scaled_attention_logits, axis=-1) # (... , seq_len_q, seq_len_k)  
  
    output = tf.matmul(attention_weights, v) # (... , seq_len_q, depth_v)  
  
    return output, attention_weights
```

# Multi-Head attention mechanism

we have 512 input dimension and we have 64 as  $W$  dimensions. so we need  $512/64 = 8$  heads.

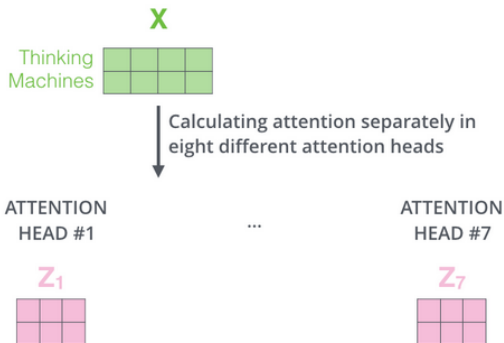
- Note: 1. we need multi head to focus on different positions. z1 contains a little bit of every other encoding, but it could be dominated by the the actual word itself.
2. expand attention layer multiple “representation subspaces”





# Multi-Head attention mechanism

model dimension = 512, number of heads = 8



# Multi-Head attention mechanism

1) Concatenate all the attention heads



2) Multiply with a weight matrix  $W^O$  that was trained jointly with the model


X



3) The result would be the  $Z$  matrix that captures information from all the attention heads. We can send this forward to the FFNN



# Multi-Head Attention



# Multi-Head-Attention

```
class MultiHeadAttention(tf.keras.layers.Layer):
    def __init__(self, d_model, num_heads):
        super(MultiHeadAttention, self).__init__()
        self.num_heads = num_heads
        self.d_model = d_model

        assert d_model % self.num_heads == 0

        self.depth = d_model // self.num_heads

        self.wq = tf.keras.layers.Dense(d_model)
        self.wk = tf.keras.layers.Dense(d_model)
        self.wv = tf.keras.layers.Dense(d_model)

        self.dense = tf.keras.layers.Dense(d_model)

    def split_heads(self, x, batch_size):
        """Split the last dimension into (num_heads, depth).
        Transpose the result such that the shape is (batch_size, n,
        """
        x = tf.reshape(x, (batch_size, -1, self.num_heads, self.depth))
        return tf.transpose(x, perm=[0, 2, 1, 3])

    def call(self, v, k, q, mask):
        batch_size = tf.shape(q)[0]

        q = self.wq(q) # (batch_size, seq_len, d_model)
        k = self.wk(k) # (batch_size, seq_len, d_model)
        v = self.wv(v) # (batch_size, seq_len, d_model)

        q = self.split_heads(q, batch_size) # (batch_size, num_heads, seq_len_q, depth)
        k = self.split_heads(k, batch_size) # (batch_size, num_heads, seq_len_k, depth)
        v = self.split_heads(v, batch_size) # (batch_size, num_heads, seq_len_v, depth)

        # scaled_attention.shape == (batch_size, num_heads, seq_len_q, depth)
        # attention_weights.shape == (batch_size, num_heads, seq_len_q, seq_len_k)
        scaled_attention, attention_weights = scaled_dot_product_attention(
            q, k, v, mask)

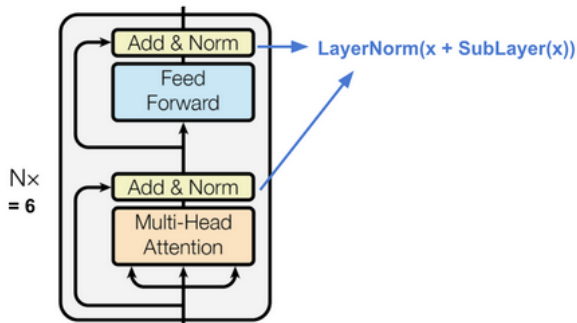
        scaled_attention = tf.transpose(scaled_attention, perm=[0, 2, 1, 3]) # (batch_size, seq_len_q, depth)

        concat_attention = tf.reshape(scaled_attention,
                                       (batch_size, -1, self.d_model)) # (batch_size, seq_len_q, d_model)

        output = self.dense(concat_attention) # (batch_size, seq_len_q, d_model)

        return output, attention_weights
```

# Encoder Layer



# Encoder Layer Tensorflow Code

```
class EncoderLayer(tf.keras.layers.Layer):
    def __init__(self, d_model, num_heads, dff, rate=0.1):
        super(EncoderLayer, self).__init__()

        self.mha = MultiHeadAttention(d_model, num_heads)
        self.ffn = point_wise_feed_forward_network(d_model, dff)

        self.layernorm1 = tf.keras.layers.LayerNormalization(epsilon=1e-6)
        self.layernorm2 = tf.keras.layers.LayerNormalization(epsilon=1e-6)

        self.dropout1 = tf.keras.layers.Dropout(rate)
        self.dropout2 = tf.keras.layers.Dropout(rate)

    def call(self, x, training, mask):

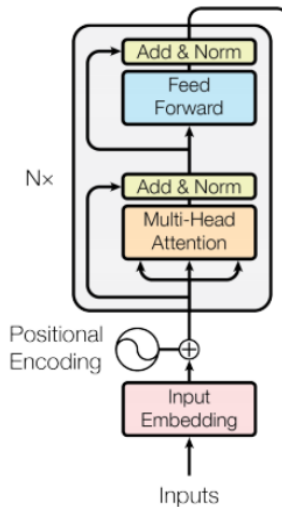
        attn_output, _ = self.mha(x, x, x, mask) # (batch_size, input_seq_len, d_model)
        attn_output = self.dropout1(attn_output, training=training)
        out1 = self.layernorm1(x + attn_output) # (batch_size, input_seq_len, d_model)

        ffn_output = self.ffn(out1) # (batch_size, input_seq_len, d_model)
        ffn_output = self.dropout2(ffn_output, training=training)
        out2 = self.layernorm2(out1 + ffn_output) # (batch_size, input_seq_len, d_model)

        return out2
```

```
def point_wise_feed_forward_network(d_model, dff):
    return tf.keras.Sequential([
        tf.keras.layers.Dense(dff, activation='relu'), # (batch_size, seq_len, dff)
        tf.keras.layers.Dense(d_model) # (batch_size, seq_len, d_model)
    ])
```

# Encoder Architecture



# Encoder Tensorflow implementation

```
class Encoder(tf.keras.layers.Layer):
    def __init__(self, num_layers, d_model, num_heads, dff, input_vocab_size,
                  maximum_position_encoding, rate=0.1):
        super(Encoder, self).__init__()

        self.d_model = d_model
        self.num_layers = num_layers

        self.embedding = tf.keras.layers.Embedding(input_vocab_size, d_model)
        self.pos_encoding = positional_encoding(maximum_position_encoding,
                                                self.d_model)

        self.enc_layers = [EncoderLayer(d_model, num_heads, dff, rate)
                            for _ in range(num_layers)]

        self.dropout = tf.keras.layers.Dropout(rate)

    def call(self, x, training, mask):

        seq_len = tf.shape(x)[1]

        # adding embedding and position encoding.
        x = self.embedding(x) # (batch_size, input_seq_len, d_model)
        x *= tf.math.sqrt(tf.cast(self.d_model, tf.float32))
        x += self.pos_encoding[:, :seq_len, :]

        x = self.dropout(x, training=training)

        for i in range(self.num_layers):
            x = self.enc_layers[i](x, training, mask)

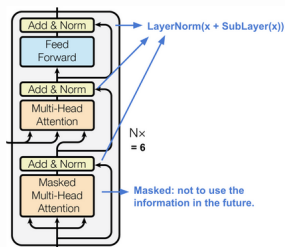
        return x # (batch_size, input_seq_len, d_model)
```



# Decoder Layer

In the decoder, the self-attention layer is only allowed to attend to earlier positions in the output sequence. This is done by masking future positions (setting them to  $-\infty$ ) before the softmax step in the self-attention calculation.

The “Encoder-Decoder Attention” layer works just like multiheaded self-attention, except it creates its Queries matrix from the layer below it, and takes the Keys and Values matrix from the output of the encoder stack.



# Decoder Layer TF code

```
class DecoderLayer(tf.keras.layers.Layer):
    def __init__(self, d_model, num_heads, dff, rate=0.1):
        super(DecoderLayer, self).__init__()

        self.mha1 = MultiHeadAttention(d_model, num_heads)
        self.mha2 = MultiHeadAttention(d_model, num_heads)

        self.ffn = point_wise_feed_forward_network(d_model, dff)

        self.layernorm1 = tf.keras.layers.LayerNormalization(epsilon=1e-6)
        self.layernorm2 = tf.keras.layers.LayerNormalization(epsilon=1e-6)
        self.layernorm3 = tf.keras.layers.LayerNormalization(epsilon=1e-6)

        self.dropout1 = tf.keras.layers.Dropout(rate)
        self.dropout2 = tf.keras.layers.Dropout(rate)
        self.dropout3 = tf.keras.layers.Dropout(rate)

    def call(self, x, enc_output, training,
             look_ahead_mask, padding_mask):
        # enc_output.shape == (batch_size, input_seq_len, d_model)


        attn1, attn_weights_block1 = self.mha1(x, x, x, look_ahead_mask) # (batch_size, target_seq_len, d_model)
        attn1 = self.dropout1(attn1, training=training)
        out1 = self.layernorm1(attn1 + x)

        attn2, attn_weights_block2 = self.mha2(
            enc_output, enc_output, out1, padding_mask) # (batch_size, target_seq_len, d_model)
        attn2 = self.dropout2(attn2, training=training)
        out2 = self.layernorm2(attn2 + out1) # (batch_size, target_seq_len, d_model)

        ffn_output = self.ffn(out2) # (batch_size, target_seq_len, d_model)
        ffn_output = self.dropout3(ffn_output, training=training)
        out3 = self.layernorm3(ffn_output + out2) # (batch_size, target_seq_len, d_model)

        return out3, attn_weights_block1, attn_weights_block2
```

# Decoder Architecture



# Decoder Architecture Tensorflow implementation

```
class Decoder(tf.keras.layers.Layer):
    def __init__(self, num_layers, d_model, num_heads, dff, target_vocab_size,
                 maximum_position_encoding, rate=0.1):
        super(Decoder, self).__init__()

        self.d_model = d_model
        self.num_layers = num_layers

        self.embedding = tf.keras.layers.Embedding(target_vocab_size, d_model)
        self.pos_encoding = positional_encoding(maximum_position_encoding, d_model)

        self.dec_layers = [DecoderLayer(d_model, num_heads, dff, rate)
                           for _ in range(num_layers)]
        self.dropout = tf.keras.layers.Dropout(rate)

    def call(self, x, enc_output, training,
            look_ahead_mask, padding_mask):

        seq_len = tf.shape(x)[1]
        attention_weights = {}

        x = self.embedding(x) # (batch_size, target_seq_len, d_model)
        x *= tf.math.sqrt(tf.cast(self.d_model, tf.float32))
        x += self.pos_encoding[:, :seq_len, :]

        x = self.dropout(x, training=training)

        for i in range(self.num_layers):
            x, block1, block2 = self.dec_layers[i](x, enc_output, training,
                                                  look_ahead_mask, padding_mask)

            attention_weights['decoder_layer{}_block1'.format(i+1)] = block1
            attention_weights['decoder_layer{}_block2'.format(i+1)] = block2

        # x.shape == (batch_size, target_seq_len, d_model)
        return x, attention_weights
```

# Transformer Decoder Output Softmax

Which word in our vocabulary  
is associated with this index?

Get the index of the cell  
with the highest value  
(**argmax**)

am

5

log\_probs



Softmax

logits

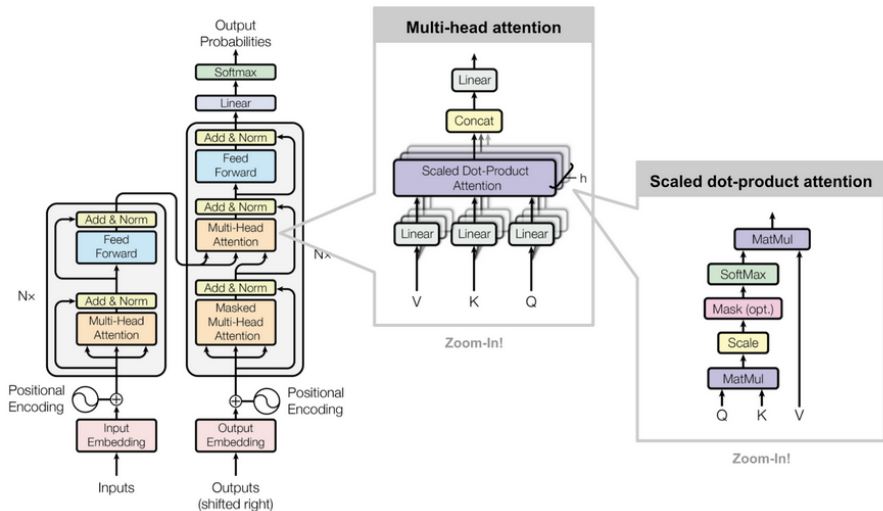


Linear

Decoder stack output



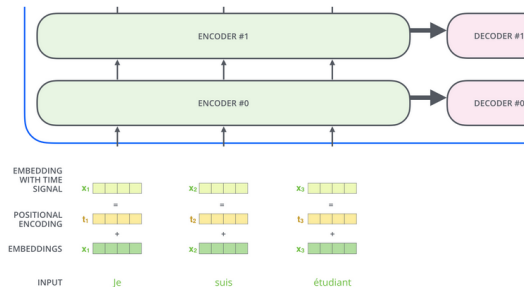
# Full Architecture



## Example

# Positional Encoding

The intuition here is that adding these values to the embedding provides meaningful distances between the embedding vectors once they're projected into Q/K/V vectors and during dot-product attention.



# Positional Encoding



# Positional Encoding

```
def get_angles(pos, i, d_model):  
    angle_rates = 1 / np.power(10000, (2 * (i//2)) / np.float32(d_model))  
    return pos * angle_rates
```

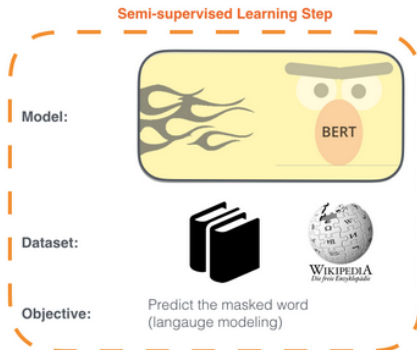
```
def positional_encoding(position, d_model):  
    angle_rads = get_angles(np.arange(position)[:, np.newaxis],  
                             np.arange(d_model)[np.newaxis, :],  
                             d_model)  
  
    # apply sin to even indices in the array; 2i  
    angle_rads[:, 0::2] = np.sin(angle_rads[:, 0::2])  
  
    # apply cos to odd indices in the array; 2i+1  
    angle_rads[:, 1::2] = np.cos(angle_rads[:, 1::2])  
  
    pos_encoding = angle_rads[np.newaxis, ...]  
  
    return tf.cast(pos_encoding, dtype=tf.float32)
```

# BERT

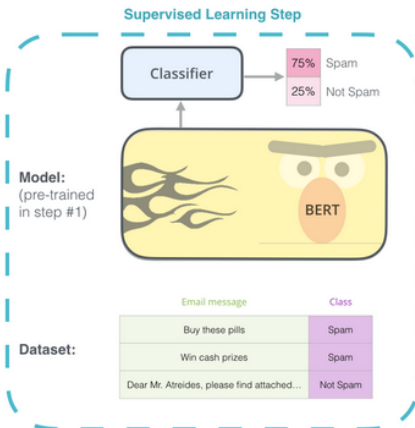
# BERT

1 - **Semi-supervised** training on large amounts of text (books, wikipedia..etc).

The model is trained on a certain task that enables it to grasp patterns in language. By the end of the training process, BERT has language-processing abilities capable of empowering many models we later need to build and train in a supervised way.



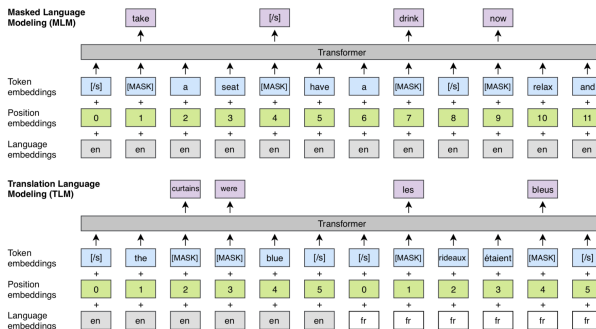
2 - **Supervised** training on a specific task with a labeled dataset.



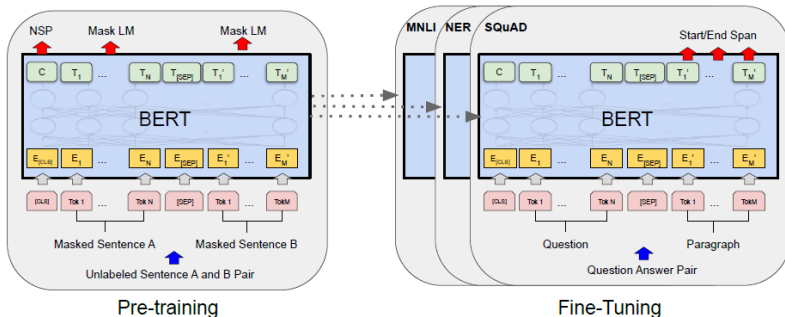
The two steps of how BERT is developed. You can download the model pre-trained in step 1 (trained on un-annotated data), and only worry about fine-tuning it for step 2. [\[Source for book icon\]](#).

# Masked Language Modeling

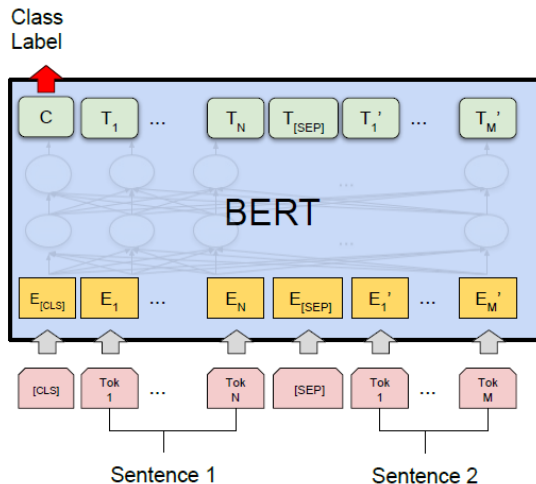
- A [CLS] token is inserted at the beginning of the first sentence and a [SEP] token is inserted at the end of each sentence.
- A sentence embedding indicating Sentence A or Sentence B is added to each token.
- A positional embedding is added to each token to indicate its position in the sequence. The concept and implementation of positional embedding are presented in the Transformer paper.



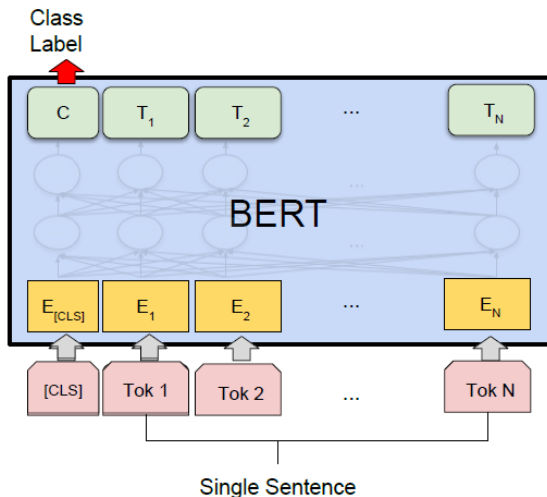
# Sentence Classification with BERT



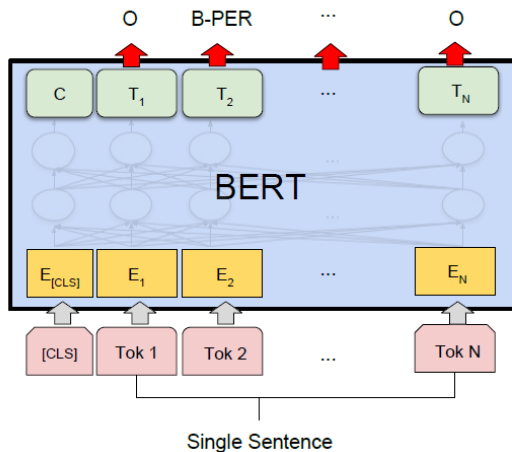
# BERT Task: Sentence pair Classification



# BERT Task: Single Sentence Classification



# BERT Task: Name Entity Recognition






# Model size and computational efficiency

## Recent trends

- Going big on model sizes - over 1 billion parameters as become the norm for SOTA



# References



- ❶ <http://jalammarm.github.io/visualizing-neural-machine-translation-mechanics-of-seq2seq-models-with-attention/>
- ❷ <https://lilianweng.github.io/lil-log/2018/06/24/attention-attention.html>
- ❸ K. Cho, B. Merriënboer, C. Gulcehre, F. Bougares, H. Schwenk, and Y. Bengio. Learning phrase representations using RNN encoder-decoder for statistical machine translation. In Arxiv preprint arXiv:1406.1078, 2014
- ❹ Bahdanau, D., Cho, K., & Bengio, Y. (2014). Neural machine translation by jointly learning to align and translate. arXiv preprint arXiv:1409.0473.
- ❺ <https://towardsdatascience.com/attention-and-its-different-forms-7fc3674d14dc>
- ❻ Luong, M.T., Pham, H. and Manning, C.D., 2015. Effective approaches to attention-based neural machine translation. arXiv preprint arXiv:1508.04025
- ❼ Devlin, J., Chang, M.W., Lee, K. and Toutanova, K., 2018. Bert: Pre-training of deep bidirectional transformers for language understanding. arXiv preprint arXiv:1810.04805.